

CS170 Operating Systems

Discussion Section

Week 2

- My office hour:
 - Wed: 4pm to 6 pm
 - Fri: 2pm to 4 pm
 - Place: CSIL
 - I will sit in the first line, near the door.
- Started coding on first project?

Project1 - Shell

Important Milestones:

1. Implement basic shell without four signs
2. Implement ">" and "<"
3. Implement basic "|"
4. Handle multiple "|"
5. Implement "&" and handle CTRL-C and D
6. Handle all 4 signs mixed together

Project1 - Shell Grading

Important Milestones:

1. Implement basic shell without four signs-30%
2. Implement ">" and "<" - 15%
3. Implement basic "|" - 20%
4. Handle multiple "|" - 15%
5. Implement "&", handle CTRL-C and D- 10%
6. Handle all 4 signs mixed together- 10%

Outline

- Implement basic shell without four signs
- Implement ">" and "<"
- Implement basic "|"
- Handle multiple "|"
- Implement "&" and handle CTRL-C and D
- Handle all 4 signs mixed together

Basic Shell

- Print a prompt

Basic Shell

- Print a prompt
- **Get the command**
 - Use fgets or gets
 - Syntax:
 gets(char *s)
 fgets(char *s, int len, FILE *fp)
 - Use "stdin" instead of FILE *fp for fgets
 - Returns 0 on "EOF"
 - Can you handle Ctrl-D using this?
 - Which one should you prefer - fgets() or gets() ?

Basic Shell

- Print a prompt
- Get a command
- Parse the command
 - Scan each character using a loop to get the tokens

Basic Shell

- Print a prompt
- Get a command
- Parse the command

Execute the command

- Use `fork()`
- Execute the command in the child process
- Use `execvp()` to take advantage of path
 - `execvp(const char *cmd, char *const argv[])`
 - What are the first and last elements of `argv`?
- Parent waits for child to terminate unless "&" is used

Something like this..

```
pid_t pID = fork();
if(pID == 0){
    ret = execvp(command_name, command_args);
}
else if(pID < 0){
    printf("Failed to fork\n");
    exit(1);
}
else{
    waitpid(-1, &status, 0);
}
```

Outline

- Implement basic shell without four signs
- Implement ">" and "<"
- Implement basic "|"
- Handle multiple "|"
- Implement "&" and handle CTRL-C and D
- Handle all 4 signs mixed together

Input and Output Redirection

- "<" means input to the command is read from file instead of stdin
- ">" means output of the command is written to file instead of stdout
- You would want to use `open()`, `dup2()`, `close()`

File Descriptors

- What are file descriptors?

File Descriptors

- File can be accessed using File descriptor
- It is non negative integer for each file
- Reserved:
 - 0(stdin), 1(stdout), 2(stderr)
- Rest of the files use $fd > 2$ when you use `open()`

Open()

- Open file:
- Defined in `unistd.h`
- Syntax
 - `open("filename",options,permissions)`
- Returns file descriptor.
 - If `fd < 0`, indicates error opening file

Open()

Example:

```
#include <fcntl.h>
```

```
#include <unistd.h>
```

```
int infilefd;
```

```
infilefd=open("infile.txt",O_RDONLY, 0);
```


Open() useful options

- O_CREAT: create if not exists
- O_RDONLY: read only
- O_WRONLY: write only
- O_TRUNC: Truncate the contents of file before writing
- O_RDWR: read and write
- Specify multiple options using "|" (or)
 - outfilefd=open("file.txt",
O_CREAT|O_WRONLY|O_TRUNC, 644);

close()

- `int close(int fd)`

Example:

```
close(outfilefd)
```

dup2()

```
#include <unistd.h>
```

```
int dup2(int fildes, int fildes2);
```

- dup2 closes the entry fildes2 of the file descriptor table, and then
- copies the pointer of entry fildes into fildes2.
- In other words it changes the pointer in: fildes2 to the pointer in fildes.

dup2(fd,fd2)




Figure 2: Before a call to `dup2 (fd, fd2)`

dup2(fd,fd2)




Figure 3: After a call to `dup2 (fd, fd2)`

Implementing "<" and ">"

- Open the file
 - Eg. `newfd=open(.....)`
- Assign newfd to 0(if "<") or 1(if ">") using `dup2`
 - Eg. `dup2(newfd,0)`
- Execute the command
 - The command reads from 0 which is your file now
- Close your file
 - Eg. `close(newfd)`
- Restore original file descriptors

Implementing "<" and ">"

- Your shell must support:
 - `cat file1`
 - `cat < file1`
 - `cat > file1`
 - `cat < file1 > file2`
 - `cat > file1 < file2`

Outline


- Implement basic shell without four signs
- Implement ">" and "<"
- Implement basic "|"
- Handle multiple "|"
- Implement "&" and handle CTRL-C and D
- Handle all 4 signs mixed together

Pipes

- Pipes are interprocess communication buffers (of course they are files).
- `#include <unistd.h>`
`int pipe(int fildes[2]);`
- On success :
 - `fildes[0]` – descriptor for read
 - `fildes[1]` – descriptor for write
 - return value 0
- On failure : return value -1

Pipes

```
int fildes[2];  
pipe(fildes);
```



Simple Example

(parent_proc | child_process)

```
int file_pipes[2];
Char some_data[] = "123";
pipe(file_pipes) ;
fork_result = fork();
...
if (fork_result == 0) { // Child
    close(file_pipes[1]);
    data = read(file_pipes[0], buffer, BUFSIZ);
    printf("Read %d bytes: %s\n", data, buffer);
    exit(EXIT_SUCCESS);
}
else { //Parent
    close(files_pipes[0]);
    data= write(file_pipes[1],some_data,strlen(some_data)+1);
    printf("Wrote %d bytes\n", data);
}
```

Implementing Pipe: ls -l | sort -n -k 5

```
int fd[2];
pid_t childpid;
pipe(fd);
if ((childpid = fork()) == 0) { /* ls is the child */
    dup2(fd[1], STDOUT_FILENO);
    close(fd[0]);
    close(fd[1]);
    execute "ls -l"
    ....
} else { /* sort is the parent */
    dup2(fd[0], STDIN_FILENO);
    close(fd[0]);
    close(fd[1]);
    execute "sort ...."
    ....
}
```

Outline

- Implement basic shell without four signs
- Implement ">" and "<"
- Implement basic "|"
- **Handle multiple "|"**
- Implement "&" and handle CTRL-C and D
- Handle all 4 signs mixed together

Handle multiple "|"

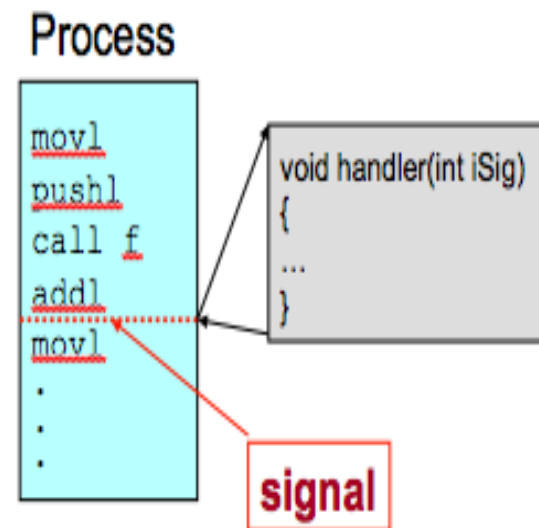
- `ls -al | grep '.c' | cut -c30- | sort -n`
- Parse the first command (everything before the first '|').
 - fork, and child will become `ls -al`
Eg. command becomes `ls -al | something`
 - Have child create `pipe()`, and `fork()`.
 - child execs `ls`.
 - grandchild processes *something* (in the same way).
- Think of a data structure!!

Outline

- Implement basic shell without four signs
- Implement ">" and "<"
- Implement basic "|"
- Handle multiple "|"
- Implement "&" and handle CTRL-C and D
- Handle all 4 signs mixed together

Signals

- **OS** communicate to an **application process** using signals
- **Signal**: A notification of an event
 - Event gains attention by OS
 - OS stops application process immediately
 - Executes signal handler completely
 - Resume process from where it left



Example



- User types Ctrl-c
 - Event gains attention of OS
 - OS stops the application process immediately, sending it a 2/SIGINT signal
 - Signal handler for 2/SIGINT signal executes to completion
 - Default signal handler for 2/SIGINT signal exits process

Signal Handling

- Every signal has a handler associated with it
 - Most default handler exit the process
- **sighandler_t signal(int iSig, sighandler_t pfHandler);**
 - Installs function **pfHandler** as the handler for signals of type **iSig**
 - **pfHandler** is a function pointer:
typedef void (*sighandler_t)(int);

Signal Handling Example

```
#include <stdio.h>
#include <signal.h>
void myHandler(int iSig) {
    printf("In myHandler with argument %d\n", iSig);
}

//invokes myhandler whenever CTRL-C is pressed
signal(SIGINT,myhandler)
```

Signal Handling options

`signal(SIGINT,SIG_IGN) //Ignore the signal`

`signal(SIGINT,SIG_DFL)`

`//Apply default signal handler(which is exit)`

Implementing '&'

1. Notice that the '&' sign can only exist as the last token, otherwise it is misplaced
2. Parent does not wait for child to complete.
3. Child sends SIGCHLD on completion

Avoid "zombies"

- When process finish execution, it has exit status to report to parent.
- So process exist in process table but has completed execution=> zombie
- Parent receives SIGCHLD when child exits
- Parent calls wait() to collect child status and removes the process

Avoid "zombies"

- Wait options:
 - WCONTINUED: The waitpid() function shall report the status of any continued child process specified by pid whose status has not been reported.
 - It suspends execution of calling thread
 - WNOHANG: The waitpid() function shall not suspend execution of the calling thread if status is not immediately available for one of the child processes specified by pid.

Handling Ctrl + D


- Use fgets and read from stdin
 - char s[100]
 - fgets(s,len(s),stdin)
 - Ctrl-D is EOF, fgets returns 0 on EOF

Outline

- Implement basic shell without four signs
- Implement ">" and "<"
- Implement basic "|"
- Handle multiple "|"
- Implement "&" and handle CTRL-C and D
- Handle all 4 signs mixed together

Putting it together!

- `ls -al | grep '.c' | cut -c30- | sort -n > foo &`



Testing

1. Compare the output by executing on terminal.

2. I will update sample input and output files on my TA page for each milestone.

<http://cs.ucsb.edu/~manish/Teaching>

3. Grading would be done using secret test files and not the ones I post.

Interesting Question

```
int main()
{
    fd = open(dir);
    printf("fd =%d",fd);
    if(fd<2)
        perror("open");
}
```

I get fd = -2 and the perror("open") prints "Success"!!!

Give a plausible explanation for this behavior.

Questions or Comments?